

● SEDIC 2026 · TRACK I · PROJECT OVERWATCH: COGNITIVE CEMA



A Dual-Branch Fusion CNN for Multi-Label RF Signal Classification

Who is transmitting, and are they hostile? NEXA reads raw radio samples, names every signal in the window (civilian comms, radar, frequency hopping, jamming) and tells the operator which tier it belongs to. It runs offline, in a browser, with a human in the loop.

[Open the live console](#)

[Read the brochure](#)

8

signal classes,
judged
independently

128,400

training and test
windows

3 / 3

judged classes
above the 80%
bar


0 servers

inference runs in
the browser

02 THE MISSION

One band. Friends, the military and an adversary, all at once.

In contested spectrum, civilian traffic, military emissions and deliberate interference overlap in time and frequency. A detector that can only name one signal per window misses the most important case: a jammer sitting on top of a real transmission. NEXA answers yes or no for each class separately, so it can report both.




One window, more than one answer: { QPSK, JAMMING }
a civilian link being jammed, reported as both.

03 HOW IT WORKS

Inside NEXA: the system architecture.

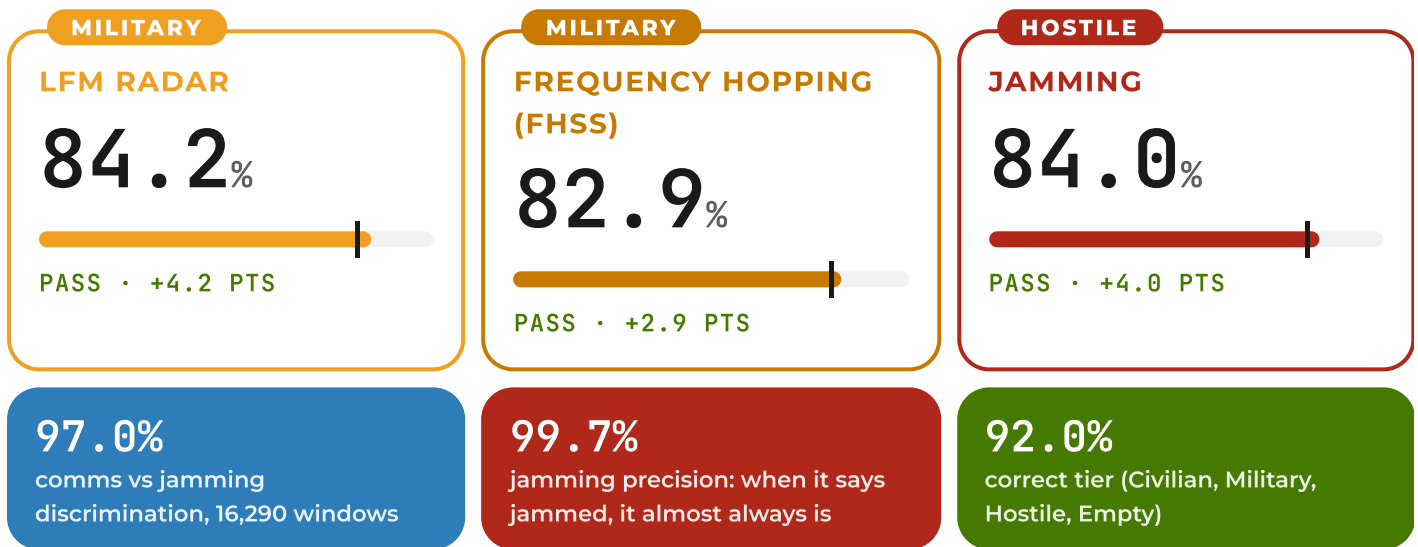
From signal acquisition to model adaptation: every capture is recorded and every decision is accountable, all on one offline workstation.



04 RESULTS

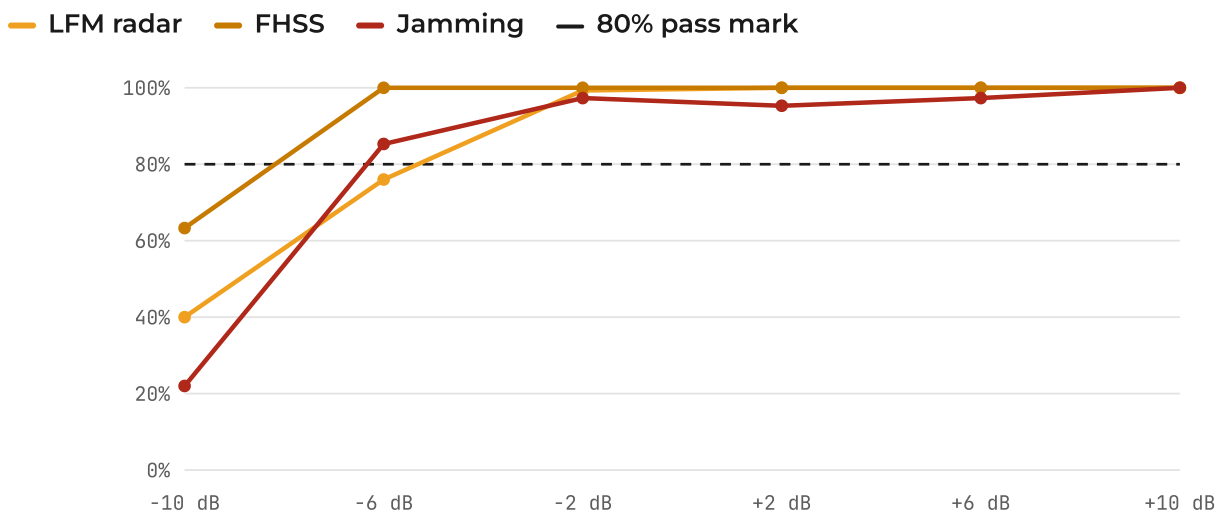
Above the bar on every judged class.

Recall on 19,260 held-out test windows, 5-model ensemble, thresholds calibrated on validation data only. The black tick on each meter is the organiser's 80% pass mark.



Recall against noise level, judged classes

Single-signal test windows. Left is noisier: at -10 dB the signal sits below the noise.

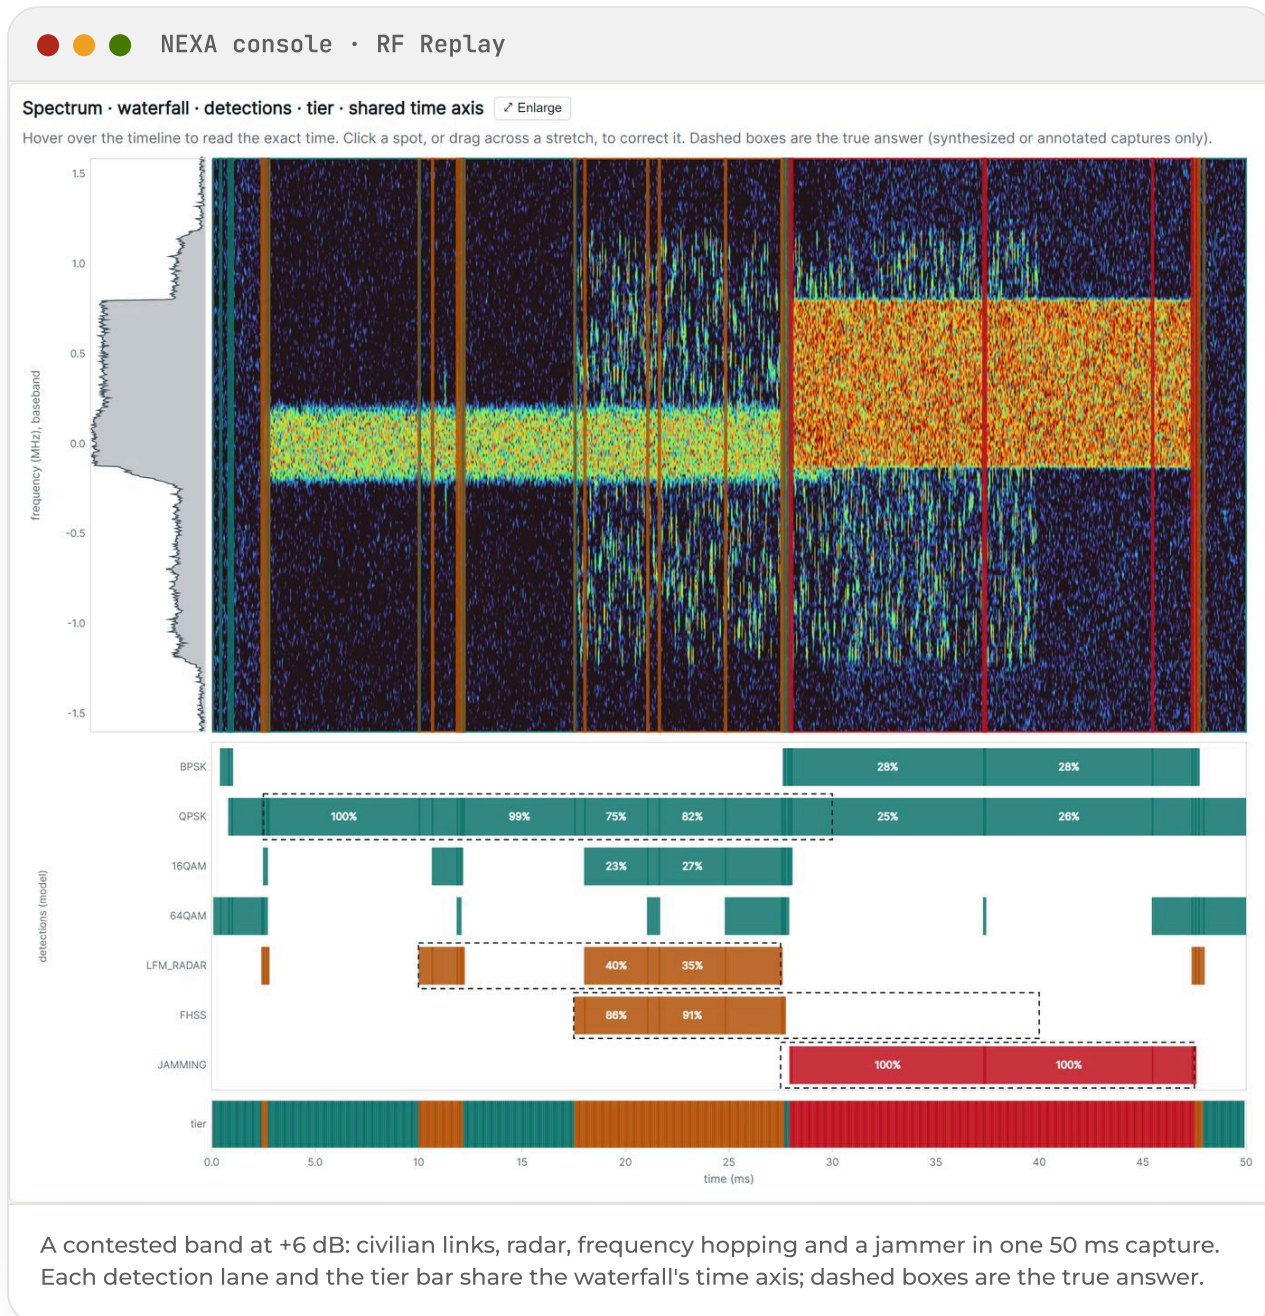


What we say before you ask

- The pass is an average over the test set. At -10 dB all three judged classes fall below 80%.
- Radar and FHSS are tuned for recall: about 6 in 10 of their alarms are correct. A missed emitter costs more than a second look.
- The drop is confined to the noisiest conditions. From -2 dB upward, radar, FHSS and jamming are each detected at 95% or better on single-signal windows.

Not just a model. A console an operator can use.

NEXA ships as a complete workstation: live view, analysis, review, reporting and retraining, with roles and an audit trail. Every screen on this page is the working console, captured as it ran.



DEMO VIDEO

Watch NEXA run

A capture goes in, the five-model ensemble classifies it in the browser, and the operator reviews the events. No server, no internet.

Watch the demo

Or [open the live console](#), or scan the codes on the back cover.


RETRAINING THAT CANNOT GET WORSE

The learning loop

NEXA learns from the operators who use it. What they correct becomes training data, but only after a second person approves it. Approved corrections fine-tune the model, mixed with the original training data so it does not forget what it already knew.

- A retrained version is offered only if it does no worse on radar, FHSS and jamming.
- Only a second person can activate it.
- Every version is kept, so any of them can be brought back.

Page: Model



RF Replay: live spectrum, waterfall and detections

Connect a receiver or replay a file as if it were live. Spectrum, waterfall, detections and the threat tier share one time axis, updated dwell by dwell. The verdict panel puts the worst tier first.

Occupancy 22.9% measured, fraction of the spectrogram above the noise floor
 Detections 32 model, grouped events, not windows
 Windows 624 hop 256 · 50.0 ms capture
 Channel 0% model, windows reported as empty spectrum

PRIMARY DETECTION
BPSK + QPSK + JAMMING
 28.00 ms · 9.36 ms long
 BPSK 28% · QPSK 25% · JAMMING 100%
 Civilian 9 · Hostile 6 · Military 17

ALL DETECTIONS - worst tier first

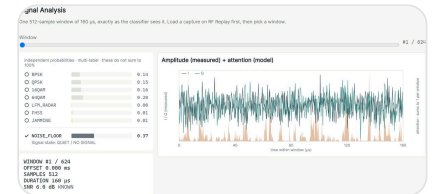
- BPSK 28% · QPSK 25% · JAMMING 100% 28.00 ms · 9.36 ms
- BPSK 28% · QPSK 26% · JAMMING 100% 37.36 ms · 8.16 ms
- BPSK 27% · QPSK 25% · 64QAM 20% · JAMMING 100% 45.44 ms · 2.00 ms
- BPSK 20% · QPSK 19% · 64QAM 16% · LFM_RADAR 45% · JAMMING 100% 47.36 ms · 0.24 ms
- BPSK 17% · QPSK 24% · 64QAM 21% · JAMMING 100% 37.28 ms · 0.16 ms

Page: RF Replay



See what the model sees

Open any 160 μ s window: measured amplitude next to the model's own attention, plus a civilian constellation showing raw IQ against recovered symbols.



Page: Signal Analysis



Assistant that cannot invent numbers

Ask in plain language. Answers come from the stored analyses; an optional local language model only rephrases the question, so it never sees or makes up data.



Page: Assistant

GOVERNANCE & ADMINISTRATION

Accountable by design



Role-based access

Operator, Analyst and Admin accounts, each with its own permissions. Every account change is logged.



Tamper-evident audit trail

Every analysis, export and deletion is logged, each entry chained to the one before, so one click proves nothing was edited.



Offline operation

The models run inside the browser through ONNX, with no internet connection and no inference server.



Four-eyes review

A correction or a missed signal becomes training data only after a different named person approves it.



Controlled reporting

PDF, Excel, CSV, JSON, a SigMF bundle or raw IQ with SHA-256 hashes, each carrying a classification banner.

06 IMPACT

What NEXA brings to contested spectrum.

Seeing who is transmitting is only useful if it leads to a faster, better decision. NEXA makes that call offline on one laptop, learns from its operators' corrections, and shows the evidence behind every figure.

OPERATIONS Clearer calls, sooner.	Threats surface first Jamming is ranked ahead of routine traffic, so attention goes where the risk is.	Overlap is reported A jammer on a live link is reported as both signals. Overlapping jamming is still found 84 to 94% of the time from -6 dB up.
	Works disconnected The full five-model ensemble runs on one laptop, with no cloud or server to depend on.	Decisions on record Two-person approval and a tamper-evident log stand behind every correction.
DEPLOYMENT Light to field, simple to run.	Runs where a browser runs Standard ONNX models and static hosting; no GPU needed to classify.	Fits existing workflows SigMF recordings in and out, with PDF, Excel, CSV and JSON reports.
	Improves in service Operator corrections feed guarded retraining, with every version kept for rollback.	Small enough to field About 182,000 parameters per model, so the five-model ensemble runs on an ordinary laptop.
RESEARCH A method that can be tested.	Evidence, labelled Each figure is marked as a requirement, a measurement, a diagnostic or an interpretation.	Honest at low SNR Recall is reported per class and per noise level, not as one average.
	Clean evaluation Thresholds are set on validation data only; every figure comes from 19,260 held-out test windows.	Reproducible Versioned configurations, fixed seeds and an automated test suite.

07 THE TEAM

Group NEXA Universiti Teknologi PETRONAS.



ET 1 Eavan Tan Radar · model and training GROUP LEADER	EY 2 Eileen Yeoh Jing Han Civilian comms · data pipeline	NT 3 Dr Noreen Talpur Universiti Teknologi PETRONAS ADVISOR	JP 4 Jessy Pang Xin Yuan Frequency hopping · evaluation	CX 5 Chua Xin Ying Jamming · inference and packaging
---	--	--	---	--



A Dual-Branch Fusion CNN for Multi-Label RF Signal Classification

AI-powered RF spectrum intelligence. Offline, in the browser,
with a human in the loop.

SCAN TO SEE IT RUN



Try the live console

Same AI, on your phone



Watch the demo

NEXA running, on video



Keep this brochure

Scan to save or share

`sedic-ai-nexa.vercel.app`

Group NEXA · Universiti Teknologi PETRONAS
SEDIC 2026 Grand Finale · 7 October 2026